

Slop Ninja

Private writing models, adversarial evaluation,
and paid inference on Bittensor

Slop Ninja Research

September 10, 2026
Draft 0.3

Abstract

Slop Ninja proposes a Bittensor subnet launching with two tasks: author and AI-origin detection (A), and author-directed text transformation (B). Miners develop private models and compete on the usefulness of their outputs. They receive network emissions for benchmark performance and can separately sell asynchronous inference at market prices denominated in the subnet’s alpha token. The protocol requires no publication of model weights, feature definitions, adapters, or training recipes.

Every customer inference request is encrypted for its assigned miner. The customer’s source text, reference writing, style profile, and editorial brief are not automatically disclosed to validators, routing services, or other miners. Customers receive encrypted results and may separately provide evidence to a specific validator of their choice. Subnet-owned benchmarks provide the authorized material that validators need to assess accuracy, preservation, style, and writing quality. Transformation miners compete against Pangram and qualified subnet AI-origin detectors. Emission eligibility requires bounded reciprocal service: detectors supply predictions and transformation miners supply authorized challenge rewrites without a per-call fee to each other. Retired benchmark rewrites feed later detection rounds. Detector reports establish observations rather than human authorship or preserved meaning.

On-chain contracts coordinate offers, commitments, deadlines, and payments; miners execute inference privately off-chain. Published hash rules determine counterpart and reviewer assignments. Every scored benchmark submission receives the required review at launch. This draft specifies the proposed architecture, training data, reward rules, and trust assumptions. It also reports the limits of the existing prototype. A reliable rewriting model, private inference service, and live reward mechanism remain to be demonstrated.

1 Purpose and scope

A useful writing model should retain a writer’s argument while allowing changes to expression. The same person may want a passage closer to their own voice, more accessible to a general audience, or more forceful without becoming less accurate. An authorship model can help describe those differences. A detector can measure one aspect of how the result is classified. Neither measurement alone establishes that the revision is good.

The initial subnet evaluates two tasks separately:

- A. Identify authors and documented human/model production histories.
- B. Transform writing toward a requested voice, optionally reducing detection by Pangram and qualified subnet AI-origin detectors, while preserving the source’s meaning, readability, and intended tone.

The joint evasion target requires Pangram’s reported AI-generated **plus AI-assisted** fraction strictly below 10%, a majority pass against the common three-detector panel, and a pass against the actual strongest qualified prior-round detector at its calibrated threshold, with preservation and quality requirements satisfied. Reaching zero

is a desirable measured outcome, not a promise of this design. A model-origin text retains that recorded origin after a successful rewrite. Detector evasion does not change its production history.

The adversarial loop connects these two tasks: detection miners learn to recognize transformed model-origin text, while transformation miners try to defeat qualified detectors from completed rounds. The benchmark fixes the opponent-assignment and scoring rules and retains the source’s production records as the models improve.

The initial scope is English prose, document-level author/origin predictions, and bounded asynchronous editing jobs. Span-level origin labels, multilingual rewards, and arbitrary-length manuscripts need their own validated datasets. There is no requirement for token streaming or interactive latency. A separate writing-improvement service is deferred beyond the initial launch; audience and editorial controls remain part of transformation briefs and preservation review.

2 Related subnets, services, and research

This representative survey covers direct task competitors and relevant evaluation and inference infrastructure. It uses project-maintained documentation and research publications checked on September 9, 2026. Subnet numbers below follow the projects’ own documentation; this is not an independent audit of live chain state. Published service claims are distinguished from results measured by this project.

2.1 Bittensor precedents

It’s AI and BitMind’s GAS supply the closest task and incentive precedents. It’s AI documents human/AI text classification on SN32, with miners serving detection requests. Its FAQ says the team does not know which models its leading miners use. This provides a precedent for evaluating detection endpoints with undisclosed implementations. GAS documents an explicit competition between detection models and generators that try to fool them [1–3].

Project	Documented approach	Relevance to Slop Ninja
It’s AI (SN32)	Human/AI text detection, including sentence-level probabilities, through miner endpoints [1,2]	Direct comparison for origin detection; compare calibration and generalization on the same independently labeled text
BitMind GAS (SN34)	Image/video/audio detector submissions and image/video generator endpoints; generation rewards include fooling detectors [3,4]	Direct precedent for adversarial incentives; text revision adds source-preservation and requested-voice constraints
Macrocosmos Finetuning (SN37)	Miners publish competition models on public Hugging Face repositories for validator evaluation [5]	Precedent for rewarding model improvements through submitted weights; Slop Ninja proposes private model serving
Chutes (SN64)	Serverless model inference and metered service, with documented confidential-serving options [6,7]	Paid inference and private serving infrastructure; writing quality requires separate task evaluation
Targon (SN4)	Managed compute and inference; confidential VMs with attestation-bound key release for proprietary workloads [8,9]	Hardware-based endpoint protection and private-model execution are established design topics

These projects establish that detection competition, adversarial generation, language-model training rewards, and paid private inference already have Bittensor precedents. Their documented mechanisms evaluate different properties. For example, GAS requires submitted detector weights, whereas Slop Ninja permits private detector services. Chutes documents confidential serving and request/response encryption, and Targon describes hardware attestation and key release. Those hardware assumptions differ from this paper’s default trust in the assigned miner after decryption. A secure execution environment alone does not establish that a revision preserved an argument.

2.2 Existing writing services

Pangram and GPTZero already offer human/model-origin classification, including mixed or assisted categories. Pangram is the required external opponent here; GPTZero is another useful comparison on rights-cleared benchmark material. Their output categories and scores must be mapped explicitly to the evaluation’s documented production labels. Agreement among detectors does not establish an author’s identity [10,11].

Grammarly documents personal voice profiles and a feature for rewriting in the user’s voice. Wordtune offers sentence and paragraph rewriting, formality changes, shortening, and expansion. QuillBot documents multiple paraphrasing modes, including custom and Humanizer modes, and warns that its Creative mode can alter meaning. These are direct comparators for author-directed revision and its editorial controls; personalization and editorial controls are existing product capabilities [12–14].

Undetectable AI explicitly markets humanization modes intended to bypass detectors. That documents an existing commercial objective; this survey does not independently establish its bypass or preservation performance [15]. Comparing such services requires the same sources, briefs, retained outputs, and blinded preservation review used for subnet miners. A product’s advertised detector score cannot substitute for that comparison. Hosted service comparisons use authorized benchmark text and do not create an exception to confidential customer delivery.

2.3 Research foundations and proposed contribution

Gero et al. model style through controls over function-word frequencies and syntactic constructions while retaining content words. This is a direct antecedent for combining lexical and grammatical features. StyleRemix uses interpretable style axes and LoRA modules for authorship obfuscation, providing a relevant comparison for controllable rewriting [16,17]. Authorship obfuscation changes evidence available to an author classifier; matching a requested writer and preserving every required claim remain separate evaluation questions.

The PAN/ELOQUENT 2024 task paired human/AI detection with generation and obfuscation intended to avoid it, offering a prior benchmark for an adversarial loop. DIPPER demonstrated paraphrase attacks on several then-current detectors and studied a retrieval-based defense. Those results motivate adaptive opponents and held-out evaluation; they do not establish performance against Pangram 4 or Slop Ninja’s future incumbents [18,19].

The proposed contribution is a text-focused protocol connecting origin-detection competition to author-directed transformation, with independent preservation, readability, and requested-voice review. Transformation miners face Pangram and qualified subnet detectors; retired, authorized rewrites supply later detector-training examples, with source families kept separate from private evaluation. Miners retain their models and can sell detection and transformation jobs through an alpha-denominated market with explicit recipient encryption. This combination requires validation against the precedents above. The survey establishes neither uniqueness across all subnets nor superior performance by an implementation that has not yet been built.

3 Participants and the two kinds of work

Miners register signed service endpoints and capabilities and retain their models. Customers buy detection and transformation jobs through separately quoted offers. Validators evaluate subnet-owned benchmark tasks and submit weights. Benchmark curators maintain permitted source material, production records, and hidden evaluation labels. A customer-facing gateway may relay ciphertext, but receives no decryption key.

Benchmark operator and coordinator denote administrative duties assigned to named curators and validators before each round: funding and managing the round, recording protocol-derived assignments, and obtaining baseline evidence. A pilot may combine these duties in one service. The task owner is the customer for private inference and the curator for a benchmark. Authorized evaluators and auditors are assigned validators; judges, readers, and authorized reviewers perform independent review for those validators. Automated judges require the qualification described in Section 8.3. The settlement adapter is the proposed contract software for customer escrow and alpha transfers, described in Section 7.

Property	Customer inference	Subnet benchmark
Purpose	Serve a customer's request	Measure miners for emissions
Launch services	A: detection; B: transformation	A: detection; B: transformation
Source of text	Customer	Authorized benchmark collection
Default input recipients	Customer and assigned miner	Curator, assigned service miners, validators, commissioned reviewers
Evaluation	Customer acceptance	Fixed scoring policy and independent review
Pangram submission	Excluded from the confidential job protocol	Required for evasion tasks
Payment	Customer's quoted fee	Performance-dependent emissions
Training reuse	No automatic reuse	Only under the collection's release terms




Figure 1: Customer inference and optional evidence disclosure. Only the addressed recipient can decrypt each envelope; endpoint handling remains a trust assumption.

Customer tasks do not become benchmark examples automatically. Validators must not obtain customer inputs through an audit exception, a shared support key, or an arbitration process. A customer may initiate a separate, explicitly addressed evidence disclosure as described in Section 6.3. This applies to both detection and transformation requests. An author's uploaded reference collection can be more identifying than the query itself and receives the same protection.

Miners may advertise specialties, prices, length limits, and delivery times. The network evaluates their observable service performance. It cannot infer how many GPU hours they used, whether they own every component, or which private weights produced a response. Emissions create an incentive to improve training; they do not reimburse unverifiable training claims.

4 Model interfaces and starter architectures

The two launch interfaces, A for detection and B for transformation, define compatible jobs and outputs. Miners may change their internal architectures, grammatical representations, search methods, and training data without changing those interfaces. Sharing a foundation model between tasks is permitted. Publishing a winning model is optional.

4.1 Author and origin detection

An author task supplies a query and reference works for a gallery of candidate authors. The response contains probabilities over the supplied author IDs. The starter compares full sparse word/grammar profiles; a learned text encoder with an author embedding head is the next architectural comparison. Reference prototypes support new customer authors without fitting a permanent classifier for each identity. Unknown-author rejection requires a separate calibrated output and withheld-author evaluation.

An origin task returns probabilities over documented production histories: human-only, model-only, and mixed human/model work. The production record includes revision order and assistance type. A model-to-model rewrite remains model-only under this initial label scheme. Pangram supplies a comparison, not the training label. The proposed encoder uses a separate origin head; author similarity must not be substituted for provenance.

ModernBERT-base [20], a 149M encoder, is a concrete starter candidate. Selection must compare it with the existing metric on new authors, topics, and registers. The subnet does not require that checkpoint or assume it will win.

4.2 Author transformation and detector evasion

Inputs are the source, independent target-author samples, requested style changes, and a preservation brief. The task declares author matching, evasion against Pangram and the common subnet AI-origin panel, or both objectives. The miner returns exact revised text or an abstention. Opponent selection and scoring are fixed by the benchmark, independently of the submitting miner. Author-directed revision of one's own or licensed writing and authorized detector-robustness research motivate this mode. A detector pass does not replace the documented production history or any disclosure required for the intended use.

The proposed starter is an 8B decoder, initially Qwen3-8B [21], with a style adapter and a candidate-selection process. Named word and grammar deviations can condition the first version. A learned profile prefix is a later comparison, contingent on better results for unseen authors. Deterministic Rust edit rules can propose small changes; the decoder can propose broader restructuring.

Train first on reviewed source/reference/brief/revision tuples. Then learn preferences among candidates that pass preservation review. Include originals, failed rewrites, copied reference passages, and fluent meaning-breaking edits as controls. A miner's own style score may guide its search, but cannot approve its output for benchmark rewards.

Transformation briefs may specify audience, clarity, accessibility, organization, rhetorical force, and tone. A lower reading level or shorter text is useful only when the brief calls for it. Suitable text can remain unchanged. Validators use independent judges and reader audits to assess preservation, readability, and requested voice; reviewer ties and disagreements remain in the data. A miner cannot approve its own revision for rewards.

Future comparisons of 4B, 8B, and 32B transformation models should use fidelity, reader preference, and cost per accepted revision. All checkpoint licenses and revisions must be pinned before a training release. The named starters are proposals, not trained subnet products.

5 Training data and evaluation separation

The two launch tasks need three linked data tables, with different supervision:

Table	Records	Supervision
Authorship	Multiple independent works per writer	Identity, date, genre, work family, provenance strength
Production	Original writing, generations, and revision chains	Recorded workflow, parent links, generator version, assistance order
Revision preferences	Sources, briefs, candidate edits, and reviews	Preservation, target voice, editorial preference, ties

Maintain distinct training, public development, private evaluation, and customer reference collections. Split authors before selecting vocabulary or fitting weights. Within evaluation authors, use independent earlier works as references and later works as queries. Keep every revision, translation, and excerpt of one source family in the same split. Hold out generator families and prompt families for origin evaluation. Match content, length, and genre across origins so that a detector cannot win by identifying the assignment format.

Global Voices supplies a prospective pool of attributed modern prose under its default CC BY terms. Qualification must distinguish original writers, translators, shared accounts, and third-party material. PLOS articles can supply factual source material and citation constraints, while their coauthored nature makes them unsuitable as unqualified single-writer labels. Item-level rights and attribution remain part of acquisition. Global Voices policy [22], PLOS article policy [23].

Historical publication alone does not prove a human-only workflow. For stronger labels and useful editing supervision, the proposed pilot commissions 50 writers to provide twelve independent pieces each and two edits of supplied drafts. Common briefs across writers permit comparisons with similar subject matter; multiple subjects per writer test consistency beyond topic. Record collaboration and AI assistance, and obtain separate rights for training, private evaluation, and any redistribution.

If collection and review prove workable, expand to 500 writers: 300 training, 100 development, and 100 private evaluation authors. A further cohort is needed for unknown-author rejection and stronger confirmation. These are acquisition targets; the writers have not been commissioned and this paper commits no funds.

Generate model-only and mixed examples from permitted sources while retaining exact prompts, model revisions, outputs, and human editing records. Synthetic restoration pairs can start from a writer’s original, create a generic draft, and train a return toward the original. Review whether that draft still contains the required information; otherwise restoration teaches invention. Keep such weak supervision separate from real writer edits.

Rewrites that fool Pangram or a subnet panel detector are hard examples for later detection rounds. Adversarial benchmark revisions retain their original production labels. They may become later detector-training examples only after retirement from private evaluation and under agreed release terms. Preserve difficult human examples alongside them. Customer text, profiles, and outputs receive no automatic training or benchmark reuse.

The existing Blog Authorship Corpus experiments remain local noncommercial research and do not supply the commercial starter dataset. Private book and correspondence material is also excluded. Source manifests must distinguish text rights from loader-code licenses. The detailed acquisition plan [24] records the current source screening and outstanding qualification work.

6 Confidential customer inference

6.1 Assigned-miner encryption

For every customer job, the protocol gives input-decryption capability only to its assigned miner. The customer already possesses the input and may retain it. Validators, gateways, storage operators, other miners, and chain observers receive no input key or alternate recipient envelope. Models execute on the assigned miner’s infrastructure; confidential jobs may not forward plaintext to an external model API, detector, or another miner.

A miner advertises a separate X25519 encryption key in a hotkey-signed record binding the network, service, key ID, validity window, and job or offer scope. Signing and encryption keys have distinct roles. The customer verifies the miner identity and finalized key registration, reserves that miner's offer, and encrypts locally before sending any payload to the network. A gateway cannot select a substitute key after receiving plaintext.

The proposed HPKE suite uses X25519 with HKDF-SHA256 for key encapsulation, HKDF-SHA256 for key derivation, and ChaCha20Poly1305 for authenticated encryption. A customer signature authenticates the envelope and its job context; HPKE Base does not itself authenticate the sender. Bind the chain, subnet, contract, job ID, assignment generation, offer hash, recipient key ID, expiry, and ciphertext to the authenticated request. Check replay state and assignment validity before processing. These are application requirements around the HPKE primitive. RFC 9180 [25].

The encrypted payload includes the source, reference works, style profile, brief, private metadata, and a commitment opening. Public commitments use a versioned, unambiguous encoding and a fresh secret 32-byte salt. An illustrative commitment is

$$C = \text{SHA256}(\text{domain} \parallel \text{context} \parallel \text{salt} \parallel \text{length} \parallel \text{payload}). \quad (1)$$

Plain text hashes, filenames, report URLs, and descriptive task titles stay inside the envelope. Padding limits exact length leakage; public size buckets, timing, prices, and counterparties remain observable.

The result is encrypted to the customer's authenticated output key. The miner knows the result it generated, but supplies no result envelope to validators. The same rule covers probabilities, explanations, and change notes: detection results can reveal sensitive inputs too.

6.2 Assignment, recovery, and the privacy boundary

Each assignment binds one miner and one encryption-key version. Rotation cannot silently redirect an existing request. A retry with another miner requires the customer to authorize a new assignment and encrypt a new request. The protocol does not escrow recovery keys or re-encrypt inputs for a fallback miner. If the customer is unavailable, a failed job expires under its settlement terms. Reassignment cannot erase a previous miner's knowledge of text it already read.

Use short-lived, preferably job-scoped keys and an explicit retention policy. Ordinary recipient encryption does not establish that a miner deleted keys, plaintext, or logs. A compromised recipient key may expose recorded ciphertext. An assigned miner can copy plaintext after decrypting it; this design protects against unauthorized protocol recipients and observers, and assumes the chosen miner follows its data-handling terms. It does not prove secrecy against that authorized endpoint. Hardware attestation or computation over encrypted data would require a separately specified and evaluated design.

There is no automatic exception for customer support, quality disputes, subnet spot checks, or Pangram verification. A customer may separately choose to disclose material to another party. In particular, the subnet must not advertise a public Pangram report as compatible with exclusive miner access.

6.3 Customer-directed evidence disclosure

A customer can ask a specific validator to inspect a job by creating a new evidence packet encrypted to that validator's authenticated key. The customer signs a disclosure context binding the job, selected validator and key version, purpose, scope, and validity window. The packet can contain source passages, the delivered result, the miner's signed response, and relevant commitment openings. The customer selects what to send. Neither the miner nor a gateway can activate this access on the customer's behalf.

This grants the selected validator access to that packet, with no decryption capability for other jobs or undisclosed material. The validator must not forward it to other reviewers or providers without further customer direction. Expiry ends authorization for new processing; it cannot erase plaintext the recipient already obtained. Keep the packet, findings, and openings private; publish only any agreed decision or commitment needed by the protocol.

The validator checks which supplied evidence opens the original commitments. A redacted excerpt does not open a commitment to a complete document. Full openings or a separately specified selective-disclosure commitment

scheme are needed for stronger binding. Incomplete evidence can support a limited opinion, but cannot establish fidelity of the whole job.

Inspection is advisory by default. A signed decision affects escrow only if the job's original terms gave that validator or selection policy authority and both parties accepted those terms. Choosing a reviewer does not itself change payment rights or extend deadlines. This route permits voluntary inspection while preserving the default assigned-miner input boundary.

7 Paid asynchronous jobs and on-chain coordination

The launch market offers A detection and B transformation jobs. The proposed job contract uses Subtensor's EVM to coordinate offers, escrow, commitments, and settlement. Inference prices are denominated in Slop Ninja's subnet alpha; actual model execution stays on miner hardware. Subtensor EVM guide [26].

7.1 Market pricing in alpha

Miners publish signed offers for bounded jobs. Each offer specifies a total alpha price, task and service version, input/output limits, delivery deadline, available capacity, quote expiry, and settlement terms. Customers select among compatible offers using price, benchmark performance, and delivery time. A higher-quality model or faster service may command a different price. There need not be one clearing price for every kind of inference.

Miners may revise their unreserved offers as demand, queue length, compute costs, and competition change. The subnet operator sets no base price, utilization target, or global price-adjustment coefficient. Miners choose their own pricing policies; customers decide whether the offers are worth buying. Alpha is the unit of account and proposed settlement asset. The inference market determines the number of alpha units charged independently of the alpha/TAO exchange market.

An atomic reservation consumes offered capacity and locks the quoted alpha amount and job terms. Later price changes apply only to later reservations. An expired or exhausted offer cannot be filled; repricing requires a new customer choice. The reservation therefore gives both parties a known price through execution and settlement, even while other offers change.

Quote discovery uses public service descriptions and the size/deadline information the customer chooses to reveal. Source text, reference writing, and the full brief remain encrypted for the selected miner. Obtaining competing offers does not distribute the plaintext to prospective bidders. If the selected miner cannot serve the bounded request, the job follows its rejection/refund terms; it cannot silently raise the price after decrypting the input.

7.2 Reservation and delivery

Subtensor's staking-v2 interface exposes allowances and stake transfers, and contracts can hold stake through their mapped coldkeys. The proposed escrow adapter would use those operations to custody unlocked, transferable alpha on the same subnet. An allowance alone does not fund a reservation. Alpha remains associated with a subnet and hotkey; settlement must account for that position and its native 10^{-9} precision. These primitives do not supply a finished job-escrow implementation [27,28].

An offer distinguishes the service price from network and transfer fees and names the fee payer. Funding succeeds only after the credited alpha principal satisfies the quote. Release and refund must reconcile actual balances, transfer permissions, and fees; incidental stake accrual must remain separate from job principal. The adapter and its accounting policy require validation before paid launch. Neither an alpha price feed nor an assumed ERC-20 interface substitutes for that work.

The proposed lifecycle is:

1. **Offer and reserve.** The customer selects a miner's signed offer, pins its key and job terms, and escrows the quoted alpha amount for a bounded input/output.
2. **Encrypt and submit.** The customer posts the salted input commitment and encrypted-delivery hash, then delivers the envelope to that miner.

3. **Accept.** The miner decrypts, checks the opening and limits, and accepts before the acceptance deadline. Failure to accept permits a refund.
4. **Execute and deliver.** The miner runs the private service, commits the result, and makes its customer-encrypted envelope available before the delivery deadline.
5. **Acknowledge or expire.** A customer acknowledgment releases payment. If there is no timely acceptance or delivery claim, the customer may claim a refund. A delivery claim without acknowledgment follows the bounded timeout policy below.

Bind acknowledgments and state transitions to chain, contract, job ID, role, nonce, and assignment generation. Bind the alpha asset's subnet ID and amount in native units as well. Authenticate the association between a Bittensor hotkey and its settlement identity. A hash or truncated address is not an authorization method. Enforce exactly one terminal settlement and explicit finality rules. A customer, miner, or keeper must submit timeout transactions; deadlines do not execute themselves.

Customers may deliver ciphertext through authenticated asynchronous endpoints and redundant encrypted storage. Small envelopes could instead travel through on-chain calldata or events, so miners can discover and complete jobs entirely through chain messaging. Payload-size, gas, and retention measurements must precede that choice. Storing ciphertext on-chain does not hide metadata or provide permanent secrecy after key compromise.

7.3 Payment without disclosing the job

A customer can consume a useful result and refuse acknowledgment. A miner can claim delivery of useless or undecryptable bytes. Public commitments alone do not distinguish those cases. Validators who cannot read the job also cannot adjudicate its meaning or writing quality.

For the first private-job pilot, use customer acknowledgment for payment and a fixed, bounded refund timeout for unacknowledged jobs. Publish that policy in the offer before acceptance. This assigns nonpayment risk to the miner, including when a customer has already consumed a valid result. Cap pilot job values and record the loss rate. It is an explicit service tradeoff, not trustless fair exchange. Public protocol evidence can resolve duplicate or late transactions; it cannot resolve a secret editorial disagreement.

Benchmark reputation can help customers choose a service without exposing their own work. It does not prove any particular unknown-input prediction correct. The customer-directed inspection route in Section 6.3 can support an agreed dispute policy. Its effect on payment must be fixed before job acceptance; a customer-selected opinion alone cannot redirect escrow. The initial pilot retains acknowledgment and bounded timeout refunds. Stronger private payment guarantees remain a launch research question.

7.4 Revenue and costs

Customer fees pay for the purchased service. Emissions reward performance on separate benchmark assignments. Customer volume does not directly increase emission rewards, because miners can transact with themselves. Paid successes cannot dilute failures on mandatory protocol service: the measured availability gate applies in aggregate and separately to each committed interface and assigned role's mandatory work. Its zero-credit and recovery-deficit rules are specified in Section 8.2. The zero-credit effect must be validated under the chain's pay-out lag; it cannot reclaim already distributed emissions. A fixed quote avoids charging for hidden, unverifiable internal token or search counts.

Miners pay their own training, inference, private search, and required candidate Pangram submissions for benchmarks. The benchmark operator or participating validators fund source-baseline reports from an explicit round budget; those reports can be reused across competing miners under the same task/version. Validators also bear retrieval, hosting, and review costs. Subsidies, if used, must identify their funder and cap before assignments begin. This design assumes no free API-credit allocation or automatic on-chain reimbursement scheme.

Required subnet-detector calls and reciprocal training requests consume reserved capacity under the service obligation in Section 8.2; the recipient pays no per-call fee within that budget. Providers bear the compute cost as a condition of emission eligibility. This obligation covers neither Pangram's external charges nor unlimited private search. Bootstrap reference services need a separately named research funder. Confidential customer jobs cannot automatically forward inputs or outputs to Pangram or a peer miner.

8 Benchmark scoring and mechanism mapping

The initial subnet has two logical tasks and two chain mechanisms: A, author/origin detection, maps to mechanism 0; B, author-directed transformation and detector evasion, maps to mechanism 1. Writing quality and meaning preservation are requirements within B, not a separate launch reward pool. Deployment must respect the documented joint UID/mechanism capacity constraint [29].

For an initial offline tournament, give A and B equal budget shares. Within A, balance author and origin tasks equally; within B, retain separate author-style, evasion, and combined modes. These are pilot allocations, not an economically optimized split. Publish any live allocation and change policy before the scoring epoch.

8.1 Detection

For a valid probability vector p and known one-hot label y , use the multiclass Brier utility:

$$B(p, y) = 1 - \frac{1}{2} \sum_k (p_k - y_k)^2. \quad (2)$$

Average query results within source groups and authors before the declared task mixture. Compare the aggregate with a fixed public baseline, then use positive improvement divided by that baseline’s remaining headroom. If the baseline is perfect, the pool has no remaining improvement and is unallocated in the offline simulation. Clip after aggregate comparison, not separately for each lucky prediction. Raw Brier utility is unsuitable for direct payout normalization: a uniform prediction in a large author gallery can already have a high value.

Report retrieval accuracy, reciprocal rank, origin calibration, and human false-positive rates separately. Fix operating thresholds on development data. Unknown or weak production records cannot become hard labels merely because Pangram agrees with a miner.

8.2 The adversarial feedback loop

Participation includes a bounded reciprocal-service obligation. A miners provide author/origin predictions on assigned B training requests; B miners provide rewrites that follow the source and brief in assigned A training requests. The request is free at the point of use, with the provider bearing its cost in exchange for eligibility for task emissions. Expected emissions need not cover every miner’s costs. The obligation concerns delivered service, not unverifiable training effort. Independent preservation, readability, and requested-tone review are required for B. A miner cannot approve its own revision for rewards.

Use released, authorized training sources and retain the production history, brief, and preservation status with each response. A later rewrite does not acquire a human-only label by defeating a detector. Unknown assistance histories remain weak labels; a private endpoint is not itself proof of model generation. Allocate complete source families to either released training material or private evaluation. Assigned miners receive the inputs needed to perform their service. Hidden labels, private review material, and detailed evaluation feedback stay with authorized evaluators until retirement; active examples are not released into reciprocal training. Customer jobs never enter this exchange automatically. Private weights remain private, but serving outputs still permits some model extraction and endpoint retention; quotas limit exposure without preventing either.

Service budgets and deterministic pairing

Freeze a global budget, request schedule, interface qualifications, and capacity limits before each evaluation epoch. Detection and generation have separate input/output limits and capacity reservations. New UIDs and traffic do not increase the global budget, and serving volume earns no reward by itself. There are no counterparty nominations or voluntary pairing offers. Customer market quotes remain separate from this protocol-assigned work.

For each protocol-issued interaction, consume the next lifetime index i_u of the requester UID slot u . The counter is persistent for that chain and subnet, including through hotkey changes and UID reuse; each assignment also pins the current registration generation, hotkey, service version, and encryption key. A requester cannot skip indices, reset them, or issue empty calls merely to advance the counter. An expired assigned slot still consumes its index

and counts as a requester failure if the requester abandons its obligation; retransmission keeps the same index, request, and peer assignment. New or replaced registrations must qualify before earning rewards; inheriting a slot counter does not inherit its former holder's performance record.

Let R_e commit the eligible, active service roster fixed for epoch e . Rank each eligible opposite-interface peer $v \neq u$ by

$$h_{u,i,v} = \text{SHA256}(\text{encode}(D, \text{chain}, \text{netuid}, u, i_u, v)), \quad (3)$$

where D is a versioned domain and encode is a canonical, length-delimited encoding. Break hash ties by UID. The first k distinct entries form the small permitted peer set. The first peer is assigned; later entries are prescribed fallbacks following recorded failure, not a menu from which the requester picks a favorable result. The pilot must publish k , capacities, and fallback limits before work begins. If the permitted set is exhausted, retain the failure; do not hash again. Hash-derived ranking and the fixed global schedule must respect provider capacity without allowing discretionary substitutions. Process capacity reservations in a canonical slot order; a request is issued only when its prescribed service windows fit the reserved capacity. Fallback stages have predeclared deadlines and minimum response windows, so later providers cannot inherit an impossible deadline. The ticket pins R_e for eligibility without mixing it into each peer's rank: changing one roster member does not reshuffle the relative order of all remaining peers.

A ticket binds the epoch, UID/index, roster root, prescribed peers and roles, permitted task, salted input commitment, size limits, deadline, and nonce. Anyone can check the scheduling metadata and sequential consumption against the finalized ledger. The text and commitment openings remain restricted to authorized recipients. Counterparty UID visibility does not establish the origin of the text. Validate authorization before inference; duplicates return the recorded response without another charge against the quota. Counter issuance and quota consumption require one authoritative, replay-protected ledger shared by validators, not independent per-validator counters.

This schedule is predictable. A miner can inspect future assignments, and another UID may have the same owner. Registration costs, bounded work schedules, and penalties for assigned failures constrain abuse; the hash does not prove operator independence or eliminate strategic abstention. Block height defines protocol-assigned slots and deadlines. Using a requester-chosen submission block as the peer seed would permit waiting for a favorable set. A short deadline narrows that opportunity without making a public block number unpredictable.

Availability from actual requests

The default Bittensor tempo is 360 blocks, approximately 72 minutes at 12 seconds per block; the actual subnet tempo and finalized epoch boundaries govern the schedule [30]. No heartbeat is required. Use authenticated outcomes of real assigned inference requests to measure availability.

For miner u and measured epoch e , let $N_{u,e}$ be the number of valid assigned task obligations and $F_{u,e}$ the number explicitly declined or lacking the required valid submission by their deadlines. An obligation is either submitting a scheduled request or answering a valid delivered request; the ticket fixes the responsible UID and role. The emission-eligibility gate is

$$L_{u,e} = \begin{cases} 1, & N_{u,e} > 0 \text{ and } 10F_{u,e} \leq N_{u,e}, \\ 0, & \text{otherwise.} \end{cases} \quad (4)$$

More than 10% failures means zero earned emission credit in both A and B for that measured epoch. Exactly 10% passes this availability gate, while quality requirements still apply. With no assigned requests there is no availability evidence and no automatic pass. Report counts and failure rates by interface as well as in aggregate. Apply the same gate separately to mandatory protocol service for each committed interface and assigned role, and require every applicable gate to pass. Purchased traffic cannot dilute failures on mandatory work, nor can cheap detection replies dilute missed transformation obligations. Successful outgoing requests cannot dilute missed provider replies. A miner needs a positive assigned workload on each interface for which it seeks eligibility; apply a role-specific gate when the schedule assigns that role work. An answer after its deadline remains a recorded late answer. A successful fallback does not erase the preceding provider's decline or failure.

Only protocol-issued requests within the fixed input, output, capacity, and response-time limits enter N . Unsolicited traffic, duplicate retries, malformed inputs, and deadlines that cannot meet the declared service window do not count. Count a protocol-issued requester obligation even if the requester withholds its payload to avoid an unfavorable peer; do not count that undelivered payload as the provider's failure. Bind the request's measured

epoch when it is issued; every counted deadline must fall within that epoch’s closed measurement window. Requests due later belong to a later window and cannot be marked late early. A valid on-time answer must have authenticated delivery evidence. A counterparty’s accusation alone cannot establish a failure. Retain disputed records for authorized validator review under a predeclared evidence rule; do not silently drop hard cases or treat an unresolved obligation as fulfilled.

Derive the next epoch’s normal active roster from registration, qualification, and the completed prior window’s request records. A small, deterministic probation allocation of real requests lets new or recovering UIDs establish availability without a self-declared online flag. Missing evidence can block normal admission without proving misconduct. Freeze each issued interaction’s roster and ordered peers; mid-epoch declines or outages invoke its existing fallback list rather than recomputing the set. The probation budget and minimum workload coverage must be measured in the pilot and fixed before live use.

The zero-credit cutoff can remove the immediate incentive to finish an already failed epoch. For each mandatory interface and role, retain a recovery deficit alongside the epoch gate:

$$D_e = \max\{0, D_{e-1} + 10F_e - N_e\}, \quad D_0 = 0. \quad (5)$$

Only authenticated, protocol-assigned work in that same class enters this update. Normal admission and emission eligibility require zero remaining deficit as well as the current availability and quality gates. Past surplus service cannot be banked; an epoch with no work cannot erase a deficit. Probation uses bounded real requests to permit recovery. Purchased jobs and other interfaces cannot retire the deficit, and key or service-version changes within a registration do not reset it. Dropping an advertised interface or role does not cancel its accrued deficit; every outstanding class must be cleared before normal admission. Replacement registrations follow their own qualification and probation; the design cannot establish whether their owners are different.

For 100 fixed obligations, 11 failures followed by 89 successful replies leave a deficit of 10. Abandoning all 100 leaves 900. Both cases earn no credit for that epoch, but require respectively 10 or 900 subsequent successful obligations to clear the deficit if no further failures occur. Capacity limits determine how many epochs recovery takes. This rule gives continued service a future eligibility benefit; it does not prove that every miner values future emissions above its costs.

The measurement window must close early enough for the relevant weight submission and commit-reveal deadlines. Publish which payout cycle settles each measured epoch. The proposed zero-credit rule cannot retroactively erase emissions already distributed by the chain. The adapter must demonstrate its effect on both mechanisms under actual weight-reveal delays and stale-weight behavior before launch; this paper assumes no custom slashing primitive.

Common evidence for transformation rewards

Per-miner peer sets allocate reciprocal training and search access. Their outputs cannot directly determine B rewards: different miners would otherwise face different detectors, and a favorable pairing could become a private reward vote. The benchmark scheduler reserves separate service capacity for a common comparison batch. It does not try to collect a full scoring matrix by burning miners’ private interaction indices.

For each shared comparison index, deterministically select three qualified A UIDs and a B batch from the frozen registry. Include the strongest qualified origin detector from the previous completed A round in the A panel; hash-rank eligible remaining A UIDs under a separate comparison domain to select the other two and the reserve order. Then hash-rank eligible B UIDs outside that panel to fill the published batch size. Filter the deterministic A reserve ranking against every UID in the B batch and the initial panel, then freeze the remaining reserve order and capacity. A reserve cannot later score its own B entry. Insufficient reserved capacity prevents issuing that comparison; it does not permit another B draw. Qualifications, sizes, capacity, tie-breaking, and the shared counter are frozen and reproducible. Miners cannot nominate a panel, join a preferred batch, or obtain another draw by declining work. Only exact UID exclusion is asserted; different UIDs may share ownership.

Rank the strongest A service by held-out origin Brier performance, subject to the declared calibration and human false-positive limits. Author-retrieval rank does not qualify an origin service. Every B entry faces the same three detectors, source texts, briefs, task modes, source-baseline evidence, and frozen settings. Collect a complete source and candidate response matrix after B candidate commitments. This is what matched assignments

mean. The three-member pilot is provisional; its cost, coverage, and susceptibility to colluding UIDs require measurement before live rewards.

If a panel service fails, apply one predetermined replacement to the whole batch and recompute that service's source and every candidate response, or void the shared comparison if reserved capacity is exhausted. Never mix per-miner panels or discard only missing cells. Detector failure is recorded against that detector's service and cannot become B nonresponse. A B miner that fails to submit receives zero on its assigned work. A missing strongest-detector result prevents a claim of success against that detector; a degraded comparison must identify its replacement and cannot silently retain the original claim.

The median panel improvement used below limits one arbitrary outlier when fewer than half the panel UIDs collude. It does not establish that bound, prove owner independence, or make private execution verifiable. One A response cannot alone determine B's utility, and A earns quality from independent labeled evaluation, not from how harshly it scores B. The strongest detector's result is also reported separately. Legitimate detector disagreement alone is neither fraud nor a replacement trigger.

For grading requests, conceal the mapping from texts to submitting B UIDs and source/candidate roles, mix human and model controls, and use one request interface. This differs from the inspectable metadata needed for peer-access verification. The service may still recognize a text or infer its role. Required signed responses cannot be replaced by private search results. Commit the complete evidence by the fixed block deadline; hash-assigned validators review every scored submission and all required checks at launch. Bootstrap rounds use explicitly funded reference services until enough competitive services qualify, and are reported separately.

8.3 Preservation before revision rewards

Each editing task has a preservation brief and hidden review material covering claims, participants, quantities, negation, uncertainty, citations, attribution, and argument relationships. Review the entire revision. Independent judges also check readability and the requested tone. A confirmed failure sets the task's preservation gate G to zero; a confirmed pass sets it to one.

Readers compare anonymized revisions in hash-derived order without detector or miner scores. Maintain unchanged sources and qualified editor baselines. Automated judges join reward allocation only after agreement and failure modes have been measured against separate reader judgments. Instructions embedded in a candidate are evaluated text, not evaluator instructions.

An unresolved judgment cannot pass by default. The round policy must resolve it or void the affected comparison for all competing miners consistently. Miner nonresponse remains a zero on assigned work; validator infrastructure failure is a separate, recorded event. Publish coverage and missingness so that conditional accuracy cannot conceal excluded difficult cases.

8.4 Transformation and improvement utility

Let $V \in [0, 1]$ be normalized positive improvement in target-style fit over the unchanged source. Estimate it from independent author models and blinded target-author or reader comparisons, retaining uncertainty. Before using a particular combination for live rewards, freeze and validate it on development examples. Topic copying, reference copying, and meaning-breaking edits must fail that validation. No miner-defined embedding can be its own reward oracle.

For Pangram evasion, let $F_r(t)$ be the reported AI-plus-assisted fraction in required observation r for text t . The initial benchmark requires three distinct completed reports for each source x and candidate z . Improvement using the least favorable pair among those submitted observations is:

$$\Delta_P(x, z) = \min_r F_r(x) - \max_r F_r(z). \quad (6)$$

The strict target additionally requires $\max_r F_r(z) < 0.10$. Three reads of one report count as one observation. The 10% cutoff is the project's declared engineering target. Any tolerance used to compare repeated reports must be reported separately and cannot widen that success criterion. Use the declared provider version and repeat policy for both sides; report sensitivity to provider variation.

Miners may pay for additional search and report attempts before commitment. Scoring uses the committed reports, including variation in their provider results. The protocol does not require disclosure of every private attempt or a fresh validator-paid Pangram inference call. A reported pass applies only to the committed observations, not to every future measurement of that text. Selection can bias the submitted set when provider results vary. We accept this limitation to keep external Pangram costs with the submitting miner; the subnet can instead require its own services to supply post-commitment responses through the reciprocal-service obligation. Our probe found equal document fractions across fresh requests for three tested texts, with auxiliary-score variation, but no observations within two percentage points of 10%. It establishes neither general determinism nor threshold stability [31].

For panel detector j , let $H_j(t) = p_j(\text{model-only} \mid t) + p_j(\text{mixed} \mid t)$, using its required authenticated response under the frozen policy. These are document-origin probabilities; Pangram's F_r is a content fraction. Define $\Delta_j(x, z) = H_j(x) - H_j(z)$ and $\Delta_S(x, z) = \text{median}_{j=1}^3 \Delta_j(x, z)$. Retain each detector's raw responses and improvements. The pilot evasion utility is

$$E(x, z) = \begin{cases} (\Delta_P(x, z) + \Delta_S(x, z))/2, & \Delta_P(x, z) \geq 0 \text{ and } \Delta_S(x, z) \geq 0, \\ 0, & \text{otherwise.} \end{cases} \quad (7)$$

The strict joint pass requires preservation, $\max_r F_r(z) < 0.10$, a majority of panel detectors satisfying $H_j(z) < \tau_j$, and a pass against the actual strongest reference service at its own threshold. Calibrate each τ_j on independent development data at the declared human false-positive operating point and freeze it before the round. These thresholds are not automatically 0.10. Report strongest-A and panel pass rates separately; the strongest-A result is not a sole veto on incremental utility. Neither Pangram improvement nor panel improvement compensates for regression on the other aggregate axis.

Pilot transformation utilities are GV , GE , and $G(V + E)/2$ for the three modes. Combined mode also requires no regression on either requested axis. Report sub-10% success separately from incremental improvement. An already suitable source can stay unchanged without receiving an invented improvement bonus.

Validate these proposed utilities against blinded reader comparisons before live reward use. They must order useful revisions above omissions, unsupported strengthening, irrelevant padding, and needless edits.

8.5 Aggregation and failure policy

Fix assignments, task mixtures, baselines, and deadlines before submissions. Balance related examples within source groups; extra cheap submissions do not increase a miner's share. Apply the availability and zero-deficit gates to each miner's earned A and B skill for the measured epoch. Then normalize eligible positive skill within each task pool and combine using the published shares. The chain mechanism split carries the corresponding aggregate task allocation.

If no miner beats a baseline, record that pool as unallocated in the offline simulation. A live adapter must select and test a chain-valid fallback before launch. This paper does not assume that zero weights, retained emissions, a new burn destination, or custom slashing are available. Payout handling under empty pools and unresolved rounds is a required deployment decision.

9 Detector evidence and assigned audits

Pangram's API provides asynchronous tasks, analyzed text, document fractions, version information, and optional public dashboard links. The returned text may differ through preprocessing, so validators must bind evidence to the analyzed bytes rather than a preview. The local reader currently requires exact equality; a normalization policy would need its own version and review. Pangram API [32].

Our synthetic probe obtained a publicly retrievable report containing the analyzed text and scores. An independent caller retrieved it without a new paid inference. The observed web endpoint is undocumented and needs an availability, retention, and usage agreement before production dependence. We have not established a provider signature, immutable model identity, or one unique canonical result per text. Recorded probe [33].

For an evasion benchmark, a miner commits its final candidate and required report IDs before the evidence deadline. Assigned validators receive the source, candidate, report URLs, score projections, and commitment openings in separate encrypted envelopes. They retrieve the provider's record directly and compare the exact text, fractions, version, and time window. Report IDs and URLs remain out of public logs and chain payloads. Anyone who obtains a public URL may still retrieve its content; encrypting the locator protects our delivery path, not Pangram's hosting.

Pangram states that it does not train Pangram 4 on customer API data. Private miner weights prevent direct downloading, but do not prevent independent black-box study through purchased outputs or leaked examples. The subnet must not assert otherwise. Provider verification and any reuse of provider labels also need suitable service terms. Pangram model card [10].

For required subnet-detector calls, retain signed responses binding the round, service identity and settings, request nonce, text commitment, and full origin-probability vector. The coordinator obtains the common panel's complete source/candidate response matrix after B candidate commitments. Miner-selected search results cannot replace it. The coordinator relays work under committed assignments; validators can verify that it followed the allocation rule. A signature establishes a service's response, not execution of particular private weights. The ticket ledger, timely-delivery evidence, and scoring path need validation independently of Pangram's public-report reader.

9.1 Assignment and audit sequence

Reciprocal partners are determined by the lifetime interaction counter and frozen eligible roster. Benchmark batches, common detector panels, and reviewer assignments use a separate, shared comparison counter and distinct hash domains. These schedules are publicly predictable. At launch, review every scored submission and all required checks, including whole-revision preservation. There is no hidden audit sample and no application-level external beacon dependency.

For each scored submission, rank eligible validator UIDs by the hash of the versioned review domain, chain, subnet, shared comparison index, submitting UID, and reviewer UID. Use canonical encoding, UID tie-breaking, three distinct reviewers, and an ordered reserve list. Exclude the exact miner UIDs involved in that comparison. The coordinator cannot choose reviewers or obtain another draw; different UIDs can still have the same owner. Reviewer capacity must be reserved before issuing the task.

1. Freeze the eligible epoch roster from qualification and actual request outcomes, service capacities, global task budgets, index-allocation rules, deadlines, and authorized auditor roster. Commit the question pool and source Pangram baselines. Publish fixed block deadlines for submissions, complete evidence, reviewer commitments, private openings, dispute resolution, and weight submission.
2. Issue reciprocal-service tickets in the protocol schedule. Consume consecutive lifetime indices, derive their ordered peers from the pinned roster, and record responses, declines, expiries, and fixed fallbacks. Training feedback uses released, authorized material; active benchmark labels and detailed review material remain private.
3. Use the shared comparison index to determine the common A panel and B batch under the frozen policy. Exclude every B UID from A reserves before freezing replacement capacity. Commit every selected service version, threshold, and reserve order. Accept one final B candidate and selected Pangram-report commitment per assigned task. The predictable opponent schedule does not allow a second draw or a revised submission after commitment.
4. Obtain the required common panel matrix through the blinded grading interface, retaining control outcomes and service failures. Apply a justified fixed replacement to every affected source/candidate cell for the batch, or void the shared comparison. Record any missing strongest-service comparison. Finalize all evidence commitments before the fixed evidence deadline; a favorable score is never grounds to switch provider or discard an observation.
5. Deliver the required evidence and inclusion proofs for every scored submission in separately encrypted envelopes to its assigned reviewers. Enforce authenticated key records, expiry, replay state, and delivery deadlines. Auditors commit answers before seeing peers' answers, then open them privately for authorized review. Resolve disagreements before the relevant weight deadline. Withhold detailed active-test feedback from revision miners until retirement; each scoring service necessarily knows its own response.

Retain every review and disagreement. A missing review follows the fixed reserve and resolution policy; it cannot default to approval. Full coverage prevents a miner from concentrating defects in a publicly predictable unreviewed subset. It does not establish reviewer honesty or guarantee that readers detect every error. A three-UID majority is useful only under a stated bound on colluding reviewers, and stake identities alone do not establish that bound. Reduce task volume to fit the funded review budget instead of silently reducing coverage. Any later sampling policy needs a separate security argument and measured cost benefit.

Subtensor's timelocked commit-reveal protects validator weights from immediate copying. Use its supported SDK/runtime path, including any beacon dependency within that native mechanism; this proposal adds no separate beacon wait for peer allocation or auditing. Weight plaintext eventually becomes public; customer inputs must not be encrypted with that eventual-reveal mechanism. Use the recipient encryption described above and do not hard-code a CRv4 schedule. Bittensor commit-reveal [34].

10 Trust assumptions and adversarial behavior

The chain can enforce authenticated accounting and commitment consistency. Recipient encryption limits who can decrypt protocol deliveries. Counter and roster commitments permit reproducible assignment of peers, comparison batches, and reviewers. Every scored submission receives the required review at launch. Each guarantee has a narrower scope than useful, confidential writing inference.

Risk	Proposed response and remaining assumption
Customer text reaches another miner or gateway	Encrypt locally to the pinned assignment; no alternate keys or automatic reassignment
Assigned miner leaks decrypted text	Minimize retention and enforce service terms; endpoint nondisclosure is not cryptographically proved
Miner retains private weights but fabricates quality	Score benchmark outputs against hidden labels and independent reviewers
Miner copies reference content or removes difficult claims	Full preservation review, copying checks, and explicit exploit controls
Miner's private search and retry costs are unverifiable	Miners fund their own attempts; score committed reports under fixed deadlines; do not reimburse claimed private costs
Validators collude or score inconsistently	Hash-assigned overlapping assignments, independent commitments, retained disagreements; honest review remains necessary
Curator or coordinator leaks labels or selects a biased task set	Independent provenance review, overlap checks, and source-group splits; commitments constrain later changes to the frozen pool and baseline evidence. Label secrecy and honest curation remain trust assumptions
Selected subnet detector favors a transformation miner	Prior-round qualification, a common three-UID panel, signed responses and blinded controls; median robustness requires a bound on colluding UIDs, and private execution remains unproved
Pangram changes or serves conflicting records	Pin the reported interface/version, retain observations, and suspend affected comparisons under a fixed policy
Customer consumes work and withholds payment	Bounded pilot exposure and explicit miner nonpayment risk; private fair exchange remains unresolved
Miner manufactures paid demand	Keep customer revenue separate from emission scoring

The benchmark also needs private source rotation, duplicate detection, and delayed release of retired examples. A service may recognize a repeated test or outsource a benchmark response; opaque version labels do not attest execution. Changing the benchmark or judge after observing a miner’s answer would invalidate the declared comparison. Publish protocol versions and aggregate outcomes while keeping customer material and active evidence private.

Reciprocal-service tickets bound aggregate free-query exposure. The predictable hash schedule does not prevent common owners from obtaining multiple UIDs or exploiting occasional favorable assignments. It provides verifiable allocation without an ownership registry. Service failure rates need authenticated request and delivery evidence; the strict greater-than-10% rule cannot be based on a recipient’s complaint alone. Predictable assignments remain mandatory, and requester abandonment consumes both an index and its failure allowance. The permitted 10% allows selective refusal; timely dishonest answers require quality evaluation and can pass availability checks. Mandatory-service gates by interface and role prevent miners from buying cheap successes to dilute failures. A persistent recovery deficit makes abandoned work delay later eligibility even after the current epoch’s credit is lost. It cannot make an operator who plans to leave value future participation. An outage can disrupt a full scoring batch, so reserve capacity and an explicit void policy remain necessary. Epoch qualification uses actual requests and deterministic probation work, without heartbeats.

11 Existing evidence and implementation status

The current Rust representation uses word occurrences and grammatical features from spaCy annotations. It scores complete sparse distributions and learns selected coordinate corrections. It does not require an LLM to enumerate or measure deterministic edit candidates. The current author/detector models and editing rules support experiments for A detection and B transformation; a demonstrated end-to-end product remains to be built.

The table reports top-1 accuracy averaged within each author, then equally across authors, then across three seeds. Query counts describe unique inputs scored in every seed; these author-macro percentages are not pooled fractions of those counts. Evaluation protocols [35,36].

Experiment	Author-macro top-1	Scope
Blog author retrieval	56.30%	600 evaluation authors in six 100-author galleries; 1,206 queries
Global Voices transfer	73.81% combined; 77.98% lexical-only	Fourteen provisionally qualified authors; 58 queries; topic/editorial confounds
Compact document projection	32.82% Blog; 25.00% Global Voices	Development-selected linear encoder; previously evaluated test corpora
Remove unselected contributions from old scorer	40.83% Blog; 29.17% Global Voices	Conditional diagnostic retaining old pooling and calibration

The projection experiment shows why 3,557 learned coordinate weights must not be mistaken for a sufficient fixed input vocabulary. The full model also retains unselected word and grammar contributions. Neither author retrieval nor these feature ablations establish meaning-preserving editing or detector evasion. In a separate edit diagnostic, replacing *may* with *must* improved author proximity despite changing the claim. Independent preservation checks are therefore central to the design. Author projection study [35], edit counterexamples [37].

Implemented components include Rust tools for corpus collection and evaluation, portable author scoring, offline challenge and submission checks, a public Pangram report reader, and an HPKE prototype that encrypts receipts for assigned validators. The receipt prototype relies on caller-supplied trusted assignments; it is not a customer-request protocol or deployed key registry. Receipt implementation [38].

11.1 Assignment and accounting stress tests

A reproducible Rust experiment uses SHA-256 rankings for 100 synthetic miner UIDs and a separate, already-eligible 100-UID reviewer roster. It evaluates 10,000 scheduled slots for each of three fixed seeds, with fixed cartel shares from 0% to 50%. The table gives pooled counts for the 20% case. Peer rates concern cartel requesters; panel and reviewer rates concern all modeled comparisons [39].

Event	Count / opportunities	Rate
Cartel peer, mandatory first assignment	1,128 / 6,000	18.80%
Cartel peer, requester chooses among three	2,875 / 6,000	47.92%
Cartel A majority, strongest A honest	1,179 / 30,000	3.93%
Cartel A majority, strongest A colluding	10,250 / 30,000	34.17%
Cartel majority among three reviewers	3,060 / 30,000	10.20%

Counting selective requester skips against each interface’s 10% limit raised favorable partners among delivered requests to 20.89%, while leaving their absolute count unchanged. Aggregate-only accounting let 34 of 60 cartel UID-epochs pass despite exceeding an interface’s limit. These results support mandatory first assignment, persistent consumed indices, and separate service gates. The recovery calculation in Equation 5 also passes the experiment’s boundary checks; its incentive effect remains conditional on valuing future participation.

Colluders can answer every request on time and reserve dishonest answers for known favorable committees. Under the modeled majority rule, complete review coverage then fails to establish correctness. Exact without-replacement calculations lower the reviewer-majority frequency from 10.11% for three reviewers to 1.47% for nine at the same fixed 20% share, with three times the review calls. Larger committees reduce this modeled frequency without proving an honest majority.

These are assignment and accounting tests, not measured model accuracy, profitable attacks, or an equilibrium analysis. The simulation assumes fixed membership and reserved capacity; it does not model adaptive UID acquisition, shared-detector fallback, delivery disputes, private-model execution, or native payouts. Three-member panels remain a pilot choice; the results do not justify calling their decisions trustless or enabling live rewards without further evidence.

The project does not yet demonstrate a reliable sub-10% rewriting model, calibrated production-origin service, reviewed commercial training release, authenticated live benchmark ledger, customer escrow contract, or private serving endpoint. The results justify a measured pilot and constrain its design; they do not establish readiness for emissions or customer funds.

12 Development and launch criteria

First, run a local tournament with 24 rights-cleared tasks across the two launch interfaces, A detection and B transformation. Evaluate detection on known-label author and origin tasks against fixed probabilistic baselines. For transformation tasks, compare unchanged text, deterministic edits, and one general editor under the same briefs. Have three hash-assigned reviewers assess preservation and preference, including intentional omissions, uncertainty changes, and padding. Freeze the scoring policy before comparison and publish aggregate agreement, coverage, cost, and failure cases.

In parallel, use synthetic jobs and test funds to verify assigned-miner encryption, customer-only result delivery, key substitution rejection, replay prevention, expiry, acknowledgment, and exactly one settlement. Exercise miner outage, wrong keys, unacknowledged consumed results, and attempted automatic reassignment. Neither a validator nor a gateway should have a decryption path. Measure encrypted payload sizes and chain costs before enabling on-chain transport. No real customer text is needed for that pilot.

Use the writer pilot to establish acquisition and review costs. Refit public starter models only on data cleared for the intended release. Evaluate author and origin models on new sources, then train the transformation adapter after enough reviewed pairs exist. Reward scaling requires demonstrated judge reliability and preservation performance, with uncertainty reported on unseen sources.

Before a live subnet launch, resolve the supported Pangram retrieval/billing integration, key and assignment registry, replay ledger, subnet-detector qualification and response verification, protocol-derived counterpart and reserve assignments, service-ticket authorization and consumption, measured reciprocal-service quotas, delivery-dispute attribution, fixed measurement, evidence, review, and weight deadlines, and complete review coverage for every scored submission, empty-pool weight policy, and mechanism allocations. Exercise exhausted quotas, duplicate tickets, concealed common control, strategic refusal, and whole-batch fallback; verify that a detector outage cannot become a B penalty or a private request's reassignment. Validate the greater-than-10% failure gate in aggregate and on each mandatory service interface and assigned role, with exactly 10% passing and zero workload supplying no automatic eligibility. Test requester abandonment, registration changes, deterministic probation, and minimum fallback response windows. Demonstrate the zero-credit rule and persistent recovery deficits, including same-class probation and resistance to resets through key or service-version changes. Verify settlement under native weight-reveal delays and payout cycles; already distributed emissions cannot be clawed back by this proposal. Before paid customer launch, separately establish acceptable nonpayment losses, retention terms, encrypted delivery reliability, and alpha escrow behavior under finality and timeouts. Verify quote expiry, capacity reservations, exact alpha accounting, release and refund permissions, and fee handling. These decisions must be concrete protocol parameters or measured acceptance criteria before funds depend on them.

The next deliverable is a small reproducible tournament and private-job pilot, followed by revised parameters based on observed failures and reader judgments. The whitepaper remains a design draft until those results support deployment.

References

- [1] It's AI. Text detection subnet. <https://github.com/It-s-AI/llm-detection>. Accessed September 9, 2026.
- [2] It's AI. Mining and model FAQ. <https://github.com/It-s-AI/llm-detection/blob/main/docs/FAQ.md>. Accessed September 9, 2026.
- [3] BitMind. GAS: Generative Adversarial Subnet. <https://github.com/BitMind-AI/bitmind-subnet>. Accessed September 9, 2026.
- [4] BitMind. GAS generative mining guide. <https://github.com/BitMind-AI/bitmind-subnet/blob/main/docs/Generative-Mining.md>. Accessed September 9, 2026.
- [5] Macrocosmos. Finetuning Subnet: Incentive overview. <https://github.com/macrocosm-os/finetuning>. Accessed September 9, 2026.
- [6] Chutes. The Chutes starter guide. <https://chutes.ai/docs/guides/starter-guide>. Accessed September 9, 2026.
- [7] Chutes. Private AI inference: What it means and how Chutes makes it verifiable. <https://chutes.ai/news/private-ai-inference-verifiable-confidential-llm-serving>. Accessed September 9, 2026.
- [8] Targon. Infrastructure. <https://targon.com/infrastructure>. Accessed September 9, 2026.
- [9] Targon. Confidential compute technical paper. <https://targon.com/releases/intel-whitepaper>. Accessed September 9, 2026.
- [10] Pangram. Pangram 4 model card. <https://www.pangram.com/research/model-card/pangram-4>. Accessed September 9, 2026.
- [11] GPTZero. Technology and API frequently asked questions. <https://gptzero.me/technology>. Accessed September 9, 2026.
- [12] Grammarly. Introducing voice features. <https://support.grammarly.com/hc/en-us/articles/23153676821773-Introducing-voice-features>. Accessed September 9, 2026.
- [13] Wordtune. Wordtune Editor explained. <https://support.wordtune.com/en/articles/8995735-wordtune-editor-explained>. Accessed September 9, 2026.
- [14] QuillBot. Paraphraser modes. <https://help.quillbot.com/hc/en-us/articles/35854318883351-What-are-modes-in-the-Quillbot-Paraphraser-and-how-do-I-use-them>. Accessed September 9, 2026.
- [15] Undetectable AI. Humanize documents. <https://help.undetectable.ai/en/article/humanize-documents-1nehl-s2/>. Accessed September 9, 2026.
- [16] Katy Gero, Chris Kedzie, Jonathan Reeve, and Lydia Chilton. Low Level Linguistic Controls for Style Transfer and Content Preservation. In *Proceedings of the 12th International Conference on Natural Language Generation*, pages 208–218. Association for Computational Linguistics, 2019. <https://aclanthology.org/W19-8628/>.

- [17] Jillian Fisher, Skyler Hallinan, Ximing Lu, Mitchell L. Gordon, Zaid Harchaoui, and Yejin Choi. StyleRemix: Interpretable Authorship Obfuscation via Distillation and Perturbation of Style Elements. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 4172–4206. Association for Computational Linguistics, 2024. <https://aclanthology.org/2024.emnlp-main.241/>.
- [18] PAN and ELOQUENT. Voight-Kampff Generative AI Authorship Verification 2024. <https://pan.webis.de/clef24/pan24-web/generated-content-analysis.html>. Accessed September 9, 2026.
- [19] Kalpesh Krishna, Yixiao Song, Marzena Karpinska, John Wieting, and Mohit Iyyer. Paraphrasing evades detectors of AI-generated text, but retrieval is an effective defense. In *Advances in Neural Information Processing Systems*, 2023. <https://arxiv.org/abs/2303.13408>.
- [20] Answer.AI and LightOn. ModernBERT-base model card. <https://huggingface.co/answerdotai/ModernBERT-base>. Accessed September 9, 2026.
- [21] Qwen Team. Qwen3-8B model card. <https://huggingface.co/Qwen/Qwen3-8B>. Accessed September 9, 2026.
- [22] Global Voices. Global Voices attribution policy. <https://globalvoices.org/about/global-voices-attribution-policy/>. Accessed September 9, 2026.
- [23] PLOS. Licenses and copyright. <https://journals.plos.org/plosone/s/licenses-and-copyright>. Accessed September 9, 2026.
- [24] Slop Ninja. Miner training sources and acquisition plan. <https://github.com/sam0x17/slopninja/blob/77b32b3/docs/miner-training-sources.md>. Repository revision 77b32b3.
- [25] Internet Engineering Task Force. RFC 9180: Hybrid Public Key Encryption. <https://www.rfc-editor.org/rfc/rfc9180.html>. Accessed September 9, 2026.
- [26] Bittensor. Subtensor EVM guide. <https://www.bittensor.com/docs/guides/evm>. Accessed September 9, 2026.
- [27] Bittensor. Staking V2 precompile. <https://www.bittensor.com/docs/guides/evm/precompiles/staking-v2>. Accessed September 9, 2026.
- [28] Bittensor. Stake from the EVM. <https://www.bittensor.com/docs/guides/evm/stake-from-evm>. Accessed September 9, 2026.
- [29] Bittensor. Subnet mechanisms. <https://www.bittensor.com/docs/guides/subnets#mechanisms>. Accessed September 9, 2026.
- [30] Bittensor. Tempo: blocks per epoch. <https://www.bittensor.com/docs/hyperparameters/tempo>. Accessed September 9, 2026.
- [31] Slop Ninja. Pangram repeatability probe. <https://github.com/sam0x17/slopninja/blob/77b32b3/docs/pangram-repeatability.md>. Repository revision 77b32b3.
- [32] Pangram. AI detection API reference. <https://docs.pangram.com/api-reference/ai-detection>. Accessed September 9, 2026.
- [33] Slop Ninja. Pangram public-result probe. <https://github.com/sam0x17/slopninja/blob/77b32b3/docs/pangram-public-results.md>. Repository revision 77b32b3.
- [34] Bittensor. Validating: Commit-reveal. <https://www.bittensor.com/docs/guides/validating#commit-reveal>. Accessed September 9, 2026.
- [35] Slop Ninja. Author projection study. <https://github.com/sam0x17/slopninja/blob/77b32b3/docs/author-projection.md>. Repository revision 77b32b3.
- [36] Slop Ninja. Global Voices author transfer. <https://github.com/sam0x17/slopninja/blob/77b32b3/docs/global-voices-transfer.md>. Repository revision 77b32b3.
- [37] Slop Ninja. Edit utility matrix and preservation counterexamples. <https://github.com/sam0x17/slopninja/blob/77b32b3/docs/edit-utility-matrix.md>. Repository revision 77b32b3.
- [38] Slop Ninja. Encrypted receipt implementation. <https://github.com/sam0x17/slopninja/blob/77b32b3/docs/receipt-envelope.md>. Repository revision 77b32b3.
- [39] Slop Ninja Research. Assignment and fulfillment gameability. <https://github.com/sam0x17/slopninja/blob/2e5f7af/docs/assignment-gameability.md>. Repository revision 2e5f7af; reproducible Rust experiment and results.